

Ankur Parag Kulkarni

Los Angeles, CA | 213-681-7805 | ankurpar@usc.edu | [LinkedIn](#) | [GitHub](#)

EDUCATION

University of Southern California

Master of Science in Computer Science, **GPA: 3.57/4**

Los Angeles, California

AUG 2024 - MAY 2026

Thakur College of Engineering and Technology

Bachelor of Engineering in Computer Engineering, **GPA: 9.67/10**

Mumbai, Maharashtra

JUL 2020 - JUN 2024

EXPERIENCE

GovRecover

Miami, Florida

Software Engineering Intern

June 2025 - Present

- Collaborated directly with the co-founder at GovRecover, joining a 5-developer team tasked with scaling backend infrastructure to expand unclaimed property automation to Texas and other state systems
- Enhance data ingestion pipelines to handle varied data formats from multiple states while ensuring reliable, self-hosted deployments with minimal downtime
- Implemented modular ETL pipelines in TypeScript and PostgreSQL to parse and standardize fixed-width and CSV formats, enabling seamless ingestion from 4+ states. Deployed and monitored services on Linode using Coolify, establishing streamlined DevOps workflows
- Achieved 99.9% service uptime across self-hosted environments, facilitating Texas system onboarding and supporting scalable, maintainable operations as GovRecover expanded

USC Marshall School of Business

Los Angeles, California

Research Assistant

Feb 2025 - Present

- USC Marshall needed scalable housing data to support post-disaster policy decisions in LA after major fires
- Build a robust data pipeline to scrape housing data while bypassing IP blocking and handling incomplete geolocation
- Built data pipelines using rotating proxies and geospatial partitioning to scrape Zillow across 500+ ZIP codes; enriched 25,000+ fire-affected properties via Azure Maps reverse geocoding and processed 100,000+ records per dataset monthly
- Improved data accuracy by 30%, enabling visualizations that guided housing trends and recovery for LA stakeholders

Mahindra Group

Mumbai, Maharashtra

Software Engineer Intern

Mar 2024–May 2024

- Mahindra sought to digitize technician training and optimize content workflows for vehicles like the XUV700 and XUV 3XO
- Build scalable solutions for content management, immersive training, and robust QA testing across platforms
- Created a React-Node CMS with AWS S3 bulk uploads, saving 200+ manual hours quarterly. Built a Unity3D VR training platform with GPT-3 NLP, reducing technician training from 3 days to 20 hours. Enhanced QA with Cypress and Playwright, increasing bug detection by 70% and cutting rollout defects by 40%.
- Delivered a comprehensive training and content system that improved operational efficiency, reduced manual effort, and ensured reliable, cross-platform user experiences

PROJECTS

Data Storytelling Web Application | React, Vite, Chart.js, Vercel

- Architected a statically-deployable, component-driven React application with modular routing and reusable UI elements, enabling the addition of 3 data visualization projects with minimal code duplication
 - Constructed client-side routing using React Router v6 with nested layouts, route-based code splitting, and fallback error handling to support seamless navigation across project pages
 - Configured environment-aware asset resolution via Vite's base path settings and deployed to Vercel using a CI-aligned deployment branch, achieving 100% successful production builds and reducing deployment setup time by 60%
- Demo: [\[link\]](#)

Smart Email Responder Workflow | Gemini 2.0, Google Cloud Functions, Python

- Engineered automated email classification and response system using Gemini API integrated with Google Cloud Functions; accurately classifies and responds to 300+ emails daily, reducing manual triage workload by 50%
 - Crafted scalable serverless functions on Google Cloud to automate real-time email processing, generating personalized responses and scheduling 50+ meetings weekly, while maintaining 99% SLA compliance
 - Improved internal communication efficiency by automating meeting invitations and follow-up notifications based on email content, significantly boosting team productivity and reducing client coordination time by 60%
- Demo: [\[link\]](#)

Medical Chatbot Fine-Tuning | Llama-3-8B, Unsloth LoRA, PyTorch

- Fine-tuned Llama-3-8B using Unsloth LoRA on 25k medical dialogues and optimized deployment via Ollama, producing a lightweight, locally hosted chatbot suitable for real-time use in clinical environments
 - Designed a privacy-preserving inference pipeline for secure, offline handling of sensitive queries, reducing external API dependency by 100% while preserving contextual accuracy and coherence in resource-constrained environments
- Demo: [\[link\]](#)

SKILLS

- Programming Languages:** Python, JavaScript, TypeScript, C++, Java
- Frontend Technologies:** Angular, React, HTML, CSS
- Backend Technologies:** Node.js, Flask, Express.js, Django
- Testing & Automation:** Selenium, Postman, Cypress
- Storage Technologies:** SQL, NoSQL, PostgreSQL, MySQL, MongoDB, DynamoDB, Redis
- Cloud & DevOps:** AWS, Azure, Docker, Kubernetes, Grafana, Git, Coolify, Linode